

# Processes

---

# Process

- A computer program in execution on a machine is a process
- More formally:
  - A **Sequential stream of Execution** in its own **address space**


# Process Address Space

- A list of memory locations from some min (usually 0) to some max that a process can read and write.
- Contains
  - the executable program
  - program's data
  - Stack
  - Associated with a process is a set of registers e.g. PC, SP and other information to run the program.

# Process =? Program

- **Program:** series of commands (e.g. C statements, assembly commands, shell commands)

**Program**  
(e.g., executable file on disk)



# Process =? Program

## ■ Anatomy of a process

- 1) More to a process than just a program
  - program is *part* of process state
  - I run *dir*, you run *dir* – same program, different processes
- 2) Less to a process than a program
  - A program can invoke more than one process to get the job done
  - e.g. cc starts up cpp, cc1, cc2

# Process =? Prog

Process

(e.g., state in memory, registers, kernel)

Mapped Segments

DLL's

Stack

{main's state}  
{A's state}

Heap

Initialized Data



Code

```
main(){  
  A();  
  ...  
}
```

Registers, PC

→ Open files, priority, user-ID, ...

# The Operating System controls the machine



# CPU State

- CPU registers contain the current state
  1. Processor Status Word (PSW): includes bits
    - Privileged or normal
    - Outcome of the last arithmetic operation (zero, -ve, +ve, overflow, carry)
    - Which interrupts are allowed and which are not
  2. Instruction Register (IR):
    - The current instruction being executed
  3. Program Counter (PC):
    - Address of the next instruction to be executed
  4. Stack Pointer (SP):
    - the address of the current stack frame, including function's local variables and return information.
  5. General purpose registers:
    - used to store addresses and data values as directed by the compiler.



# Memory Contents

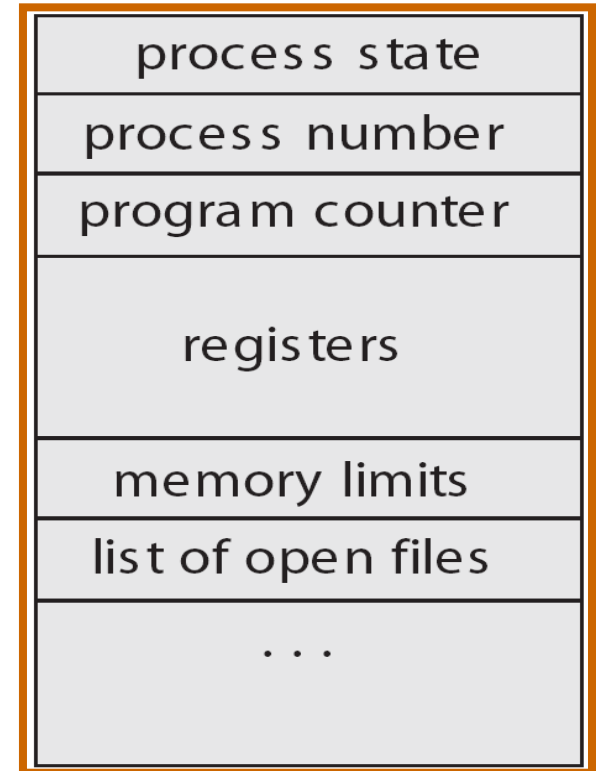
- Only a small part of an application's data can be stored in registers. The rest is in memory.
- Typically divided into a few segments:
  - Text/application code
    - read-only
    - might be shared by a number of processes?
  - Data
    - The application's predefined data structures
  - Heap
    - An area from which space can be allocated dynamically at runtime, using functions like **new** or **malloc**.
  - Stack
    - Where register values are saved
    - local variables allocated
    - Return Address of subroutine calls
    - All the addressable memory together is called?
- The process's address space.

# Environment

- Contains the relationships with other entities
- A process does not exist in a vacuum
- It typically has connections with other entities, such as
  - A terminal where the user is sitting.
  - Open files
  - Communication channels to other processes, possibly on other machines.

# Process Control Block

- The OS keeps all the data it needs about a process in the process control block (PCB)
- Thus another definition of a process:
  - “the entity described by a PCB”
- This includes many of the data items described above, or at least pointers to where they can be found
  - e.g. for the address space



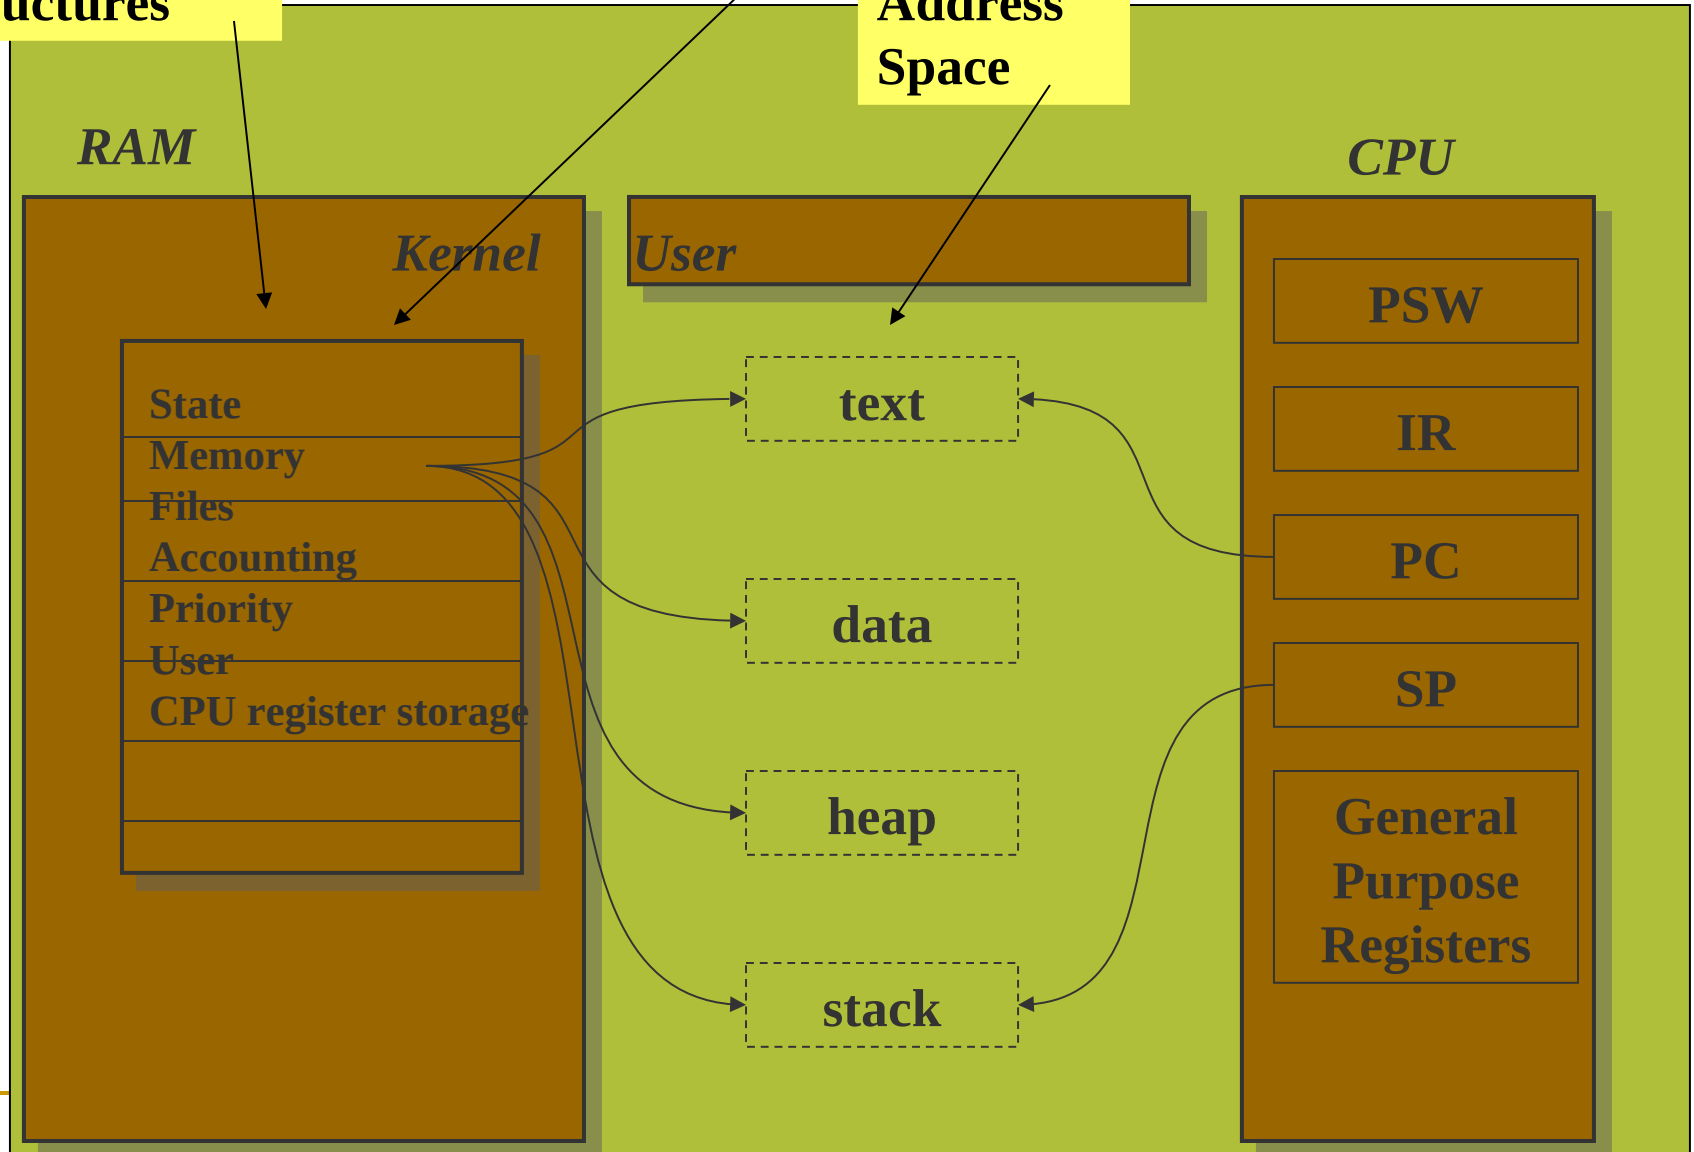
# Process Identification

- Process ID, a unique numeric identifier
- User ID
  - Who runs the process. Why?
  - Used to determine what access rights the process has

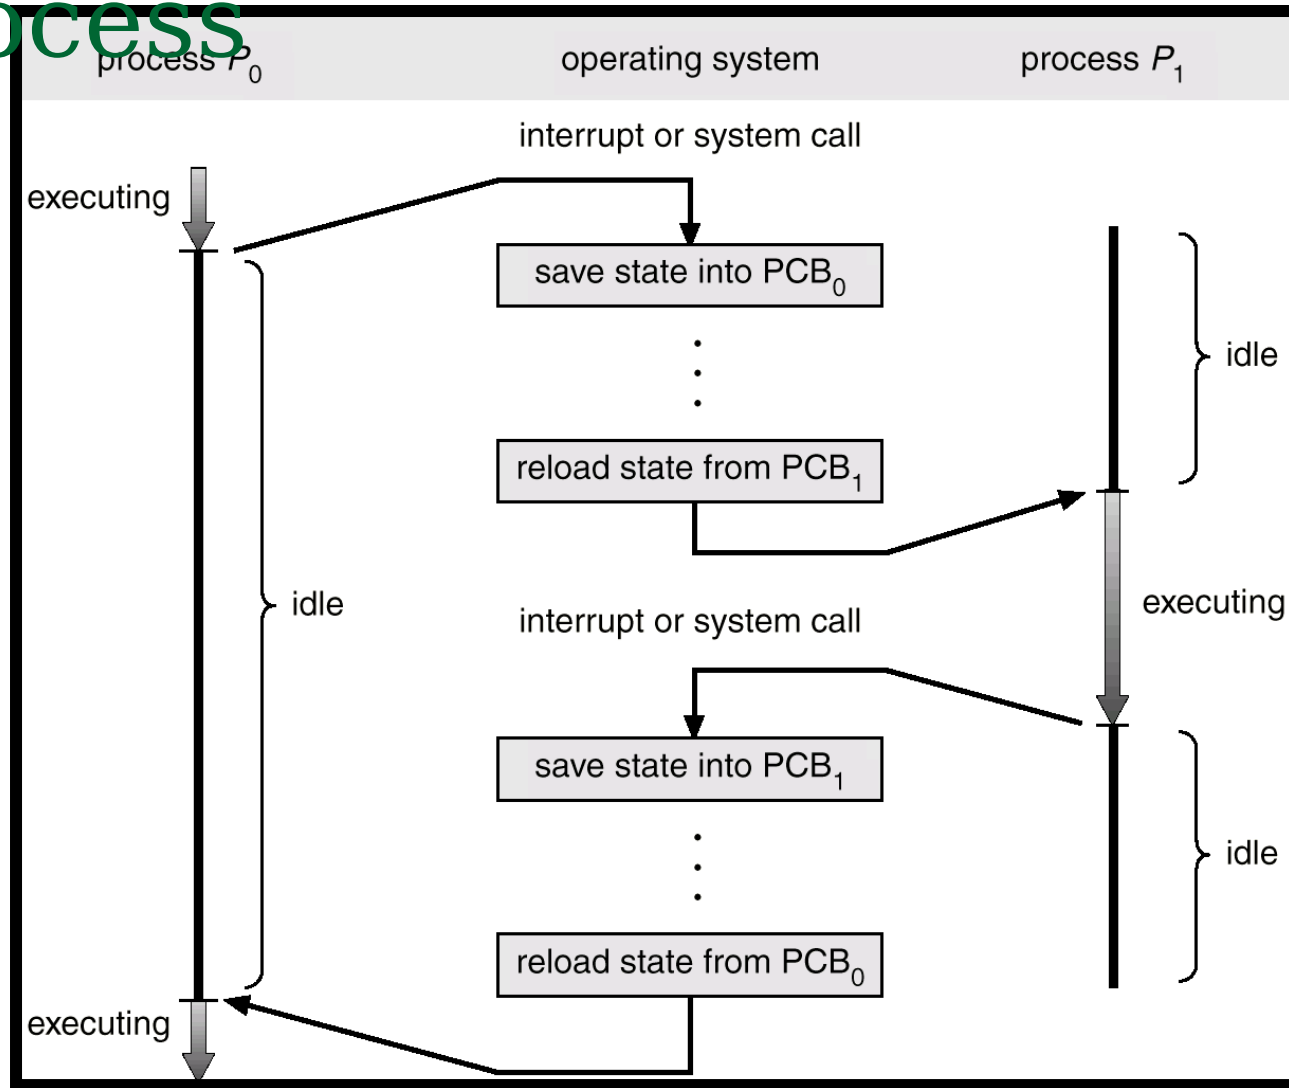
# Process Control

OS Data Structures

Process Address Space



# CPU Switch From Process to Process

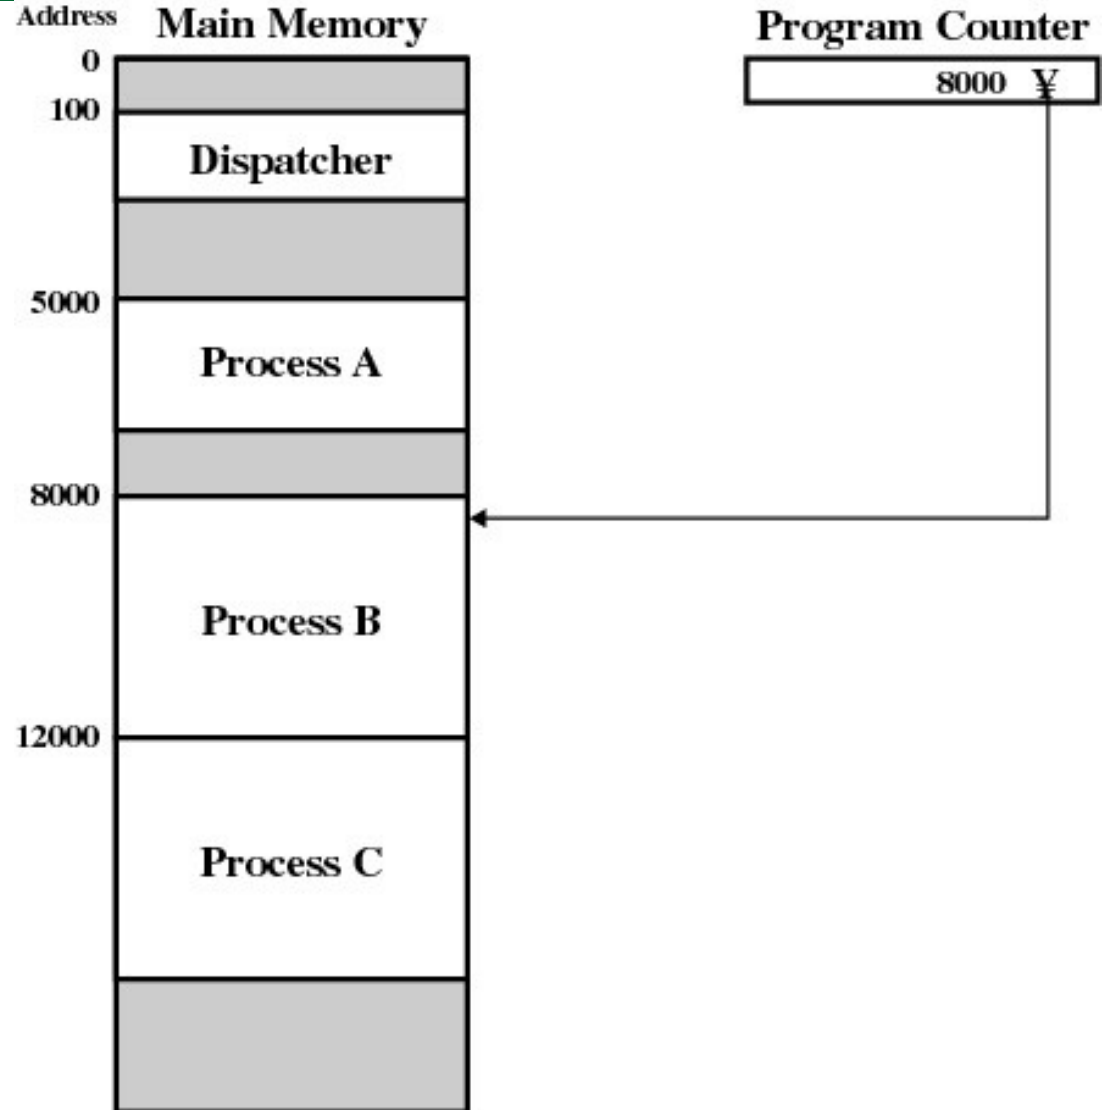


---

# CPU Switch From Process to Process

- Switching a process requires
  - Saving the state of old process
  - Loading the saved state of the new process
- This is called **Context Switch**
- Part of OS responsible for switching the processor among the processes is called **Dispatcher**

# Process Example






# Process States


- At any given time a process is either running or not running
- Number of states
  - Running
  - Not Running
- When the OS creates a process, the process is entered into which state?
  - Not Running

# Two-state process model



- Number of processes **Running** at a particular time?
- Number of processes **Not Running** at a particular time?
- Data Structure?
- Processes that are **Not Running** at a particular time should be kept in some sort of a **queue**

# Two-state process model



## Dispatcher is now redefined:

- Moves processes to the waiting queue
- Remove completed/aborted processes
- Select the next process to run

# How process state changes1

Ready


→ a := 1  
b := a + 1  
c := b + 1  
read a file  
a := b - c  
c := c \* b  
b := 0

Ready

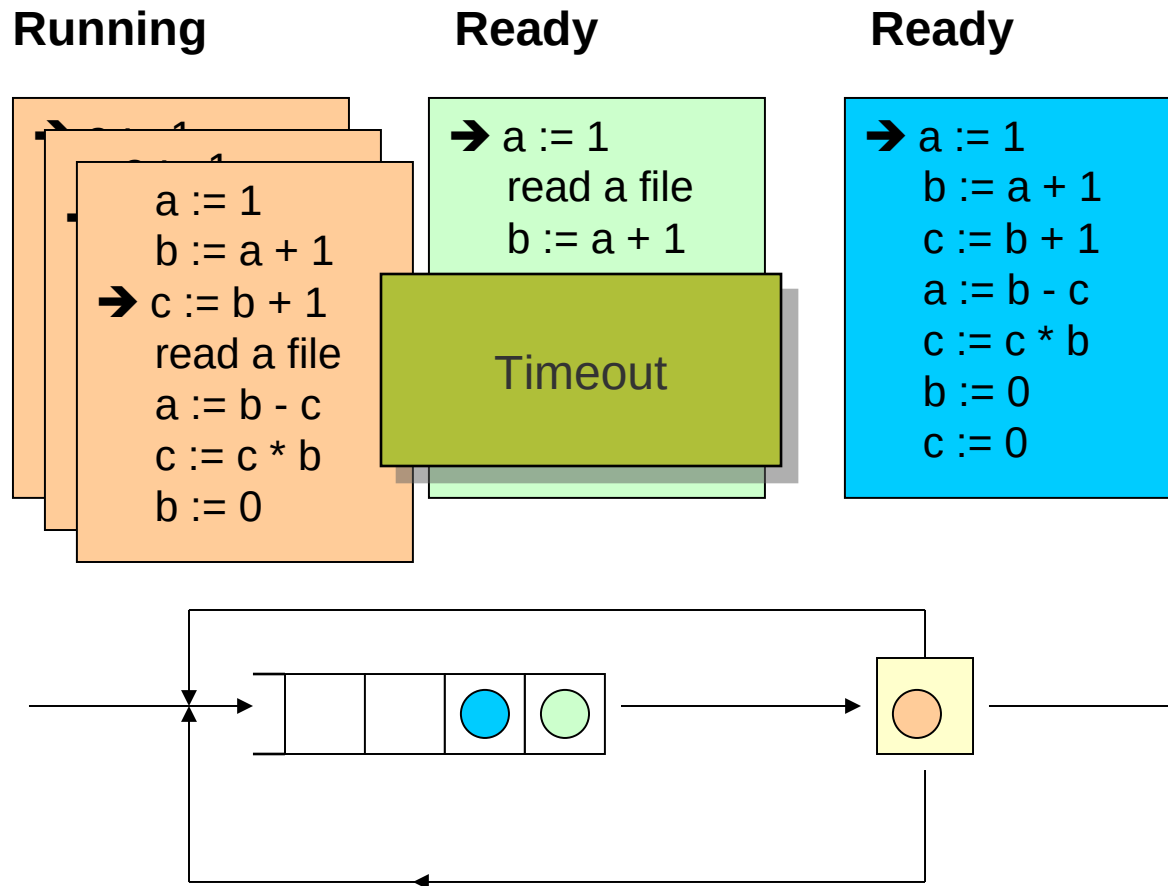
→ a := 1  
read a file  
b := a + 1  
c := b + 1  
a := b - c  
c := c \* b  
b := 0

Ready

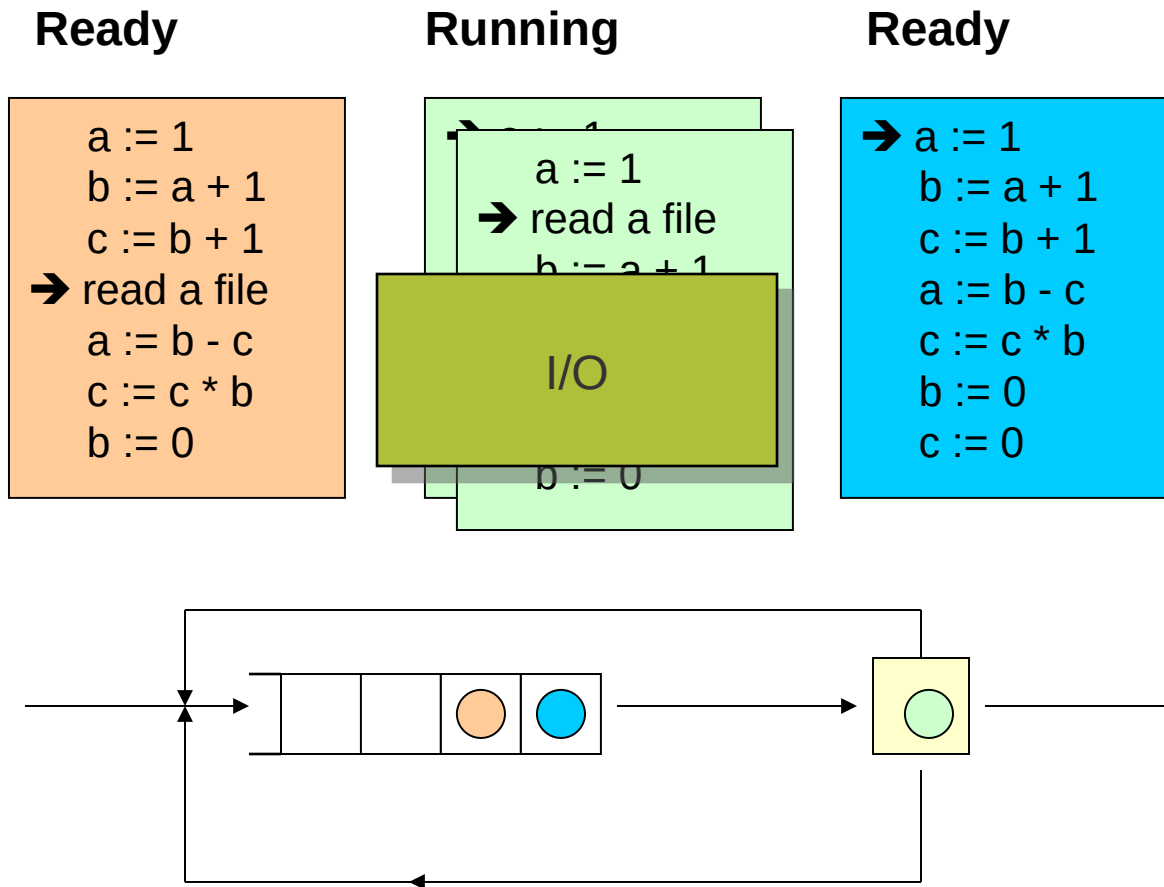
→ a := 1  
b := a + 1  
c := b + 1  
a := b - c  
c := c \* b  
b := 0  
c := 0



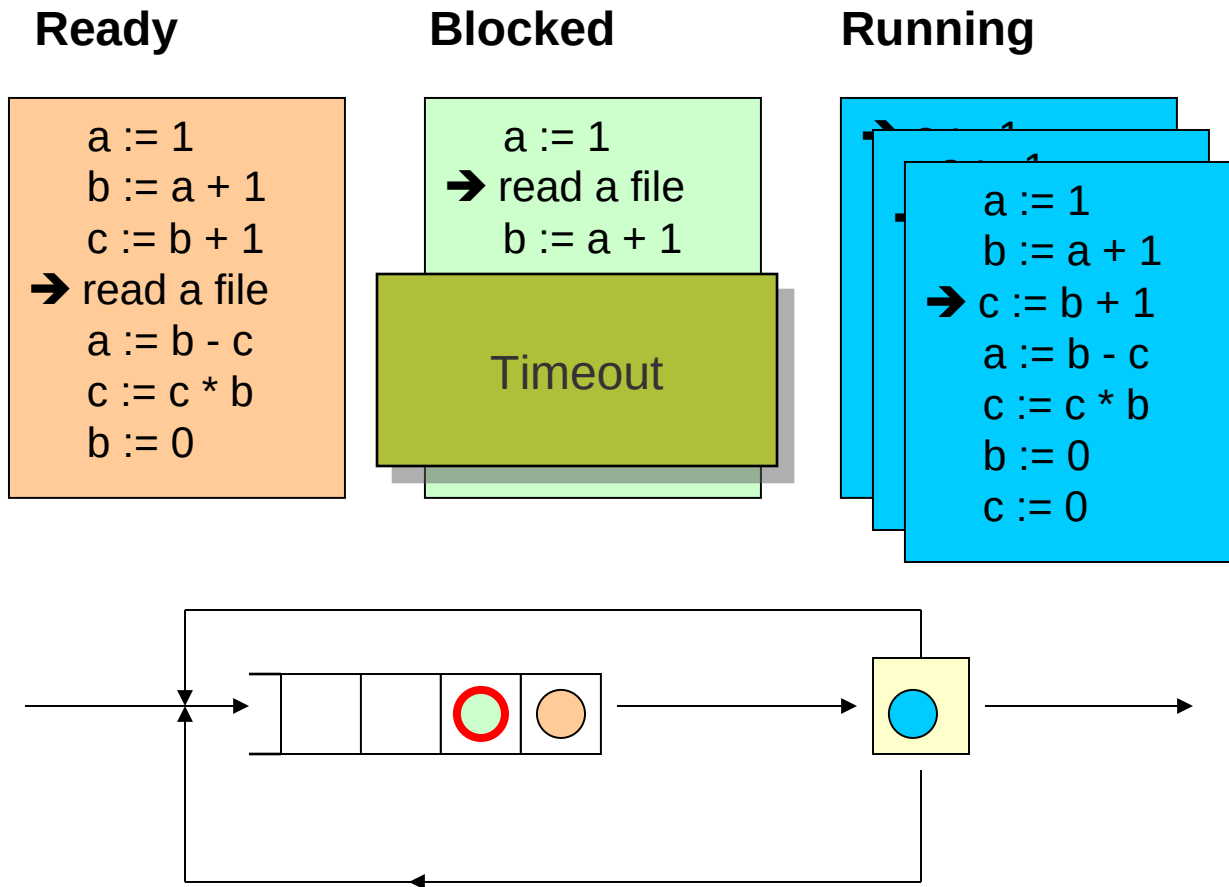
# How process state changes2



# How process state changes3



# How process state changes4



# How process state changes5

## Running

```
a := 1
b := a + 1
c := b + 1
→ read a file
a := b - c
c := c * b
b := 0
```


## Blocked

```
a := 1
→ read a file
b := a + 1
```

I/O

## Ready

```
a := 1
b := a + 1
c := b + 1
→ a := b - c
c := c * b
b := 0
c := 0
```





# How process state changes

**Blocked**


```
a := 1
b := a + 1
c := b + 1
→ read a file
a := b - c
c := c * b
b := 0
```

**Blocked**

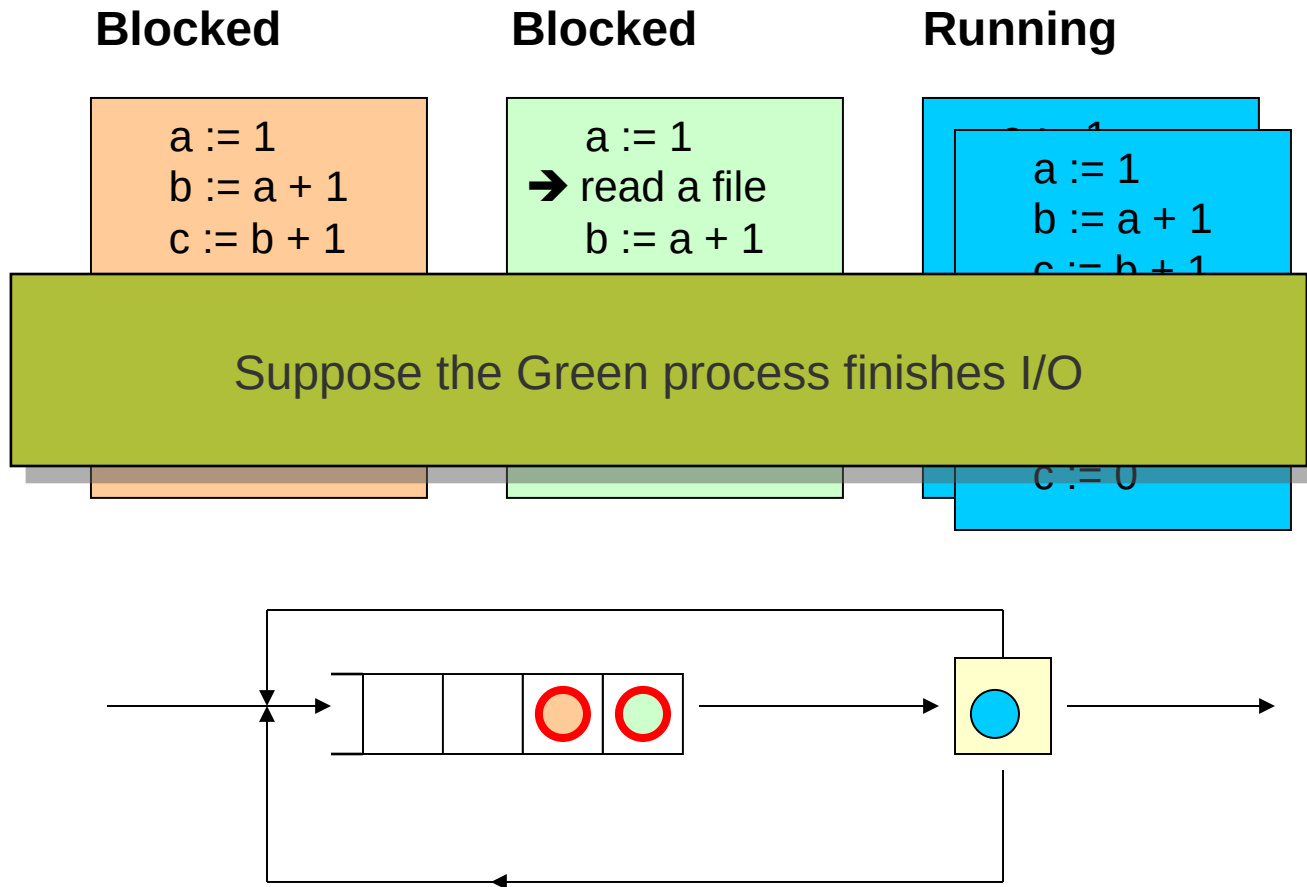
```
a := 1
→ read a file
b := a + 1
c := b + 1
a := b - c
c := c * b
b := 0
```

**The Next Process to Run cannot be simply selected from the front**

```
→ a := b - c
c := c * b
b := 0
c := 0
```



# How process state changes?



# How process state changes8

**Blocked**

```
a := 1
b := a + 1
c := b + 1
→ read a file
a := b - c
c := c * b
b := 0
```


**Ready**

```
a := 1
read a file
→ b := a + 1
```

Timeout

**Running**

```
a := 1
b := a + 1
c := b + 1
a := b - c
c := c * b
→ b := 0
c := 0
```




# How process state changes9

**Blocked**


```
a := 1
b := a + 1
c := b + 1
→ read a file
a := b - c
c := c * b
b := 0
```

**Running**



**Ready**

```
a := 1
b := a + 1
c := b + 1
a := b - c
c := c * b
b := 0
→ c := 0
```



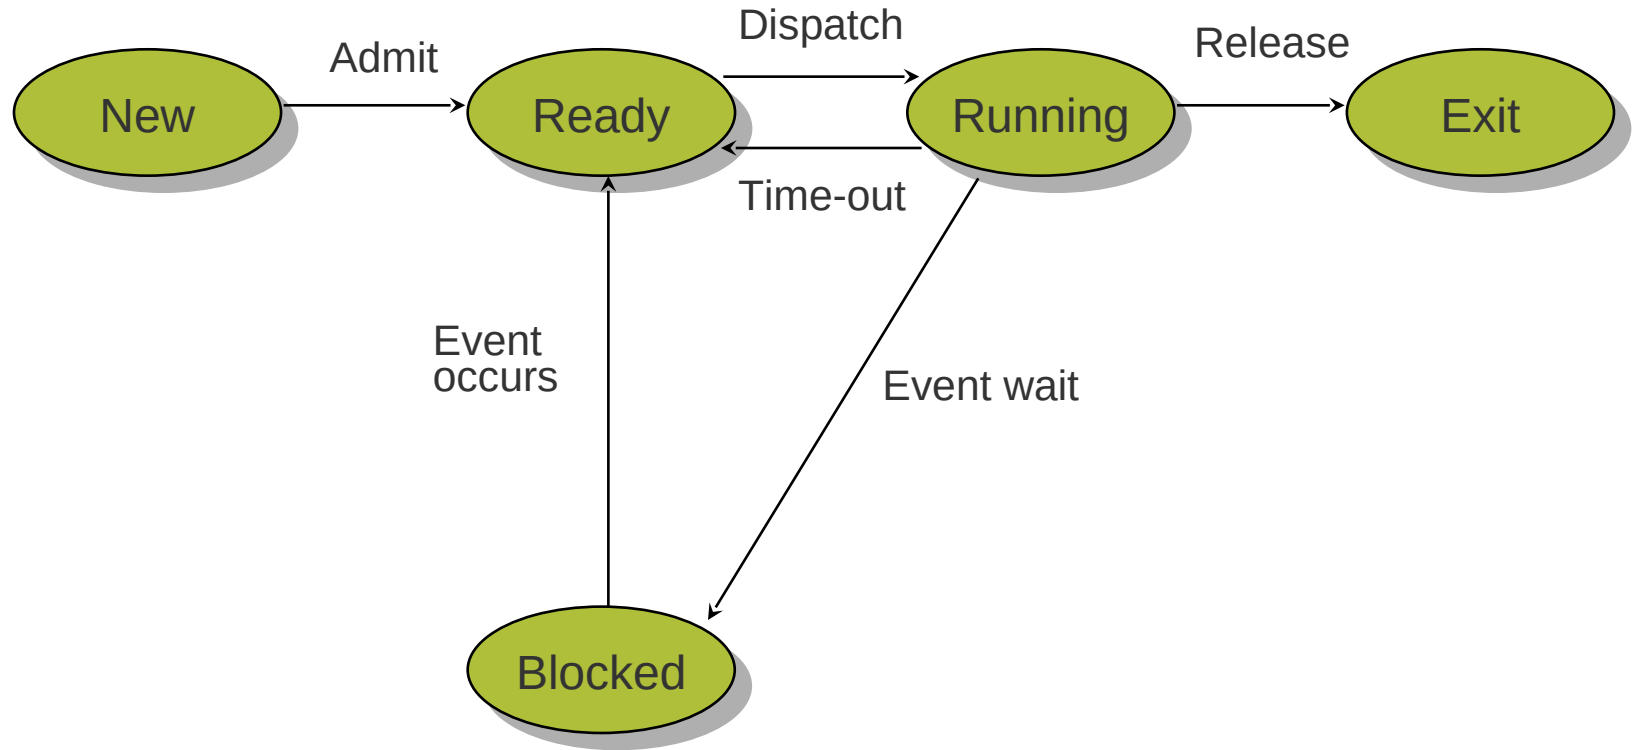
# Problem in Two-state Process model

- A process may be waiting for I/O request
- A single queue for both the ready to run and waiting processes
- The dispatcher cannot simply select the process at the front, it can be a busy process
- In the worst case, it has to scan the whole queue to find the next process to run
- Solution?
- Split the Not Running state to:
  - Waiting
  - Ready

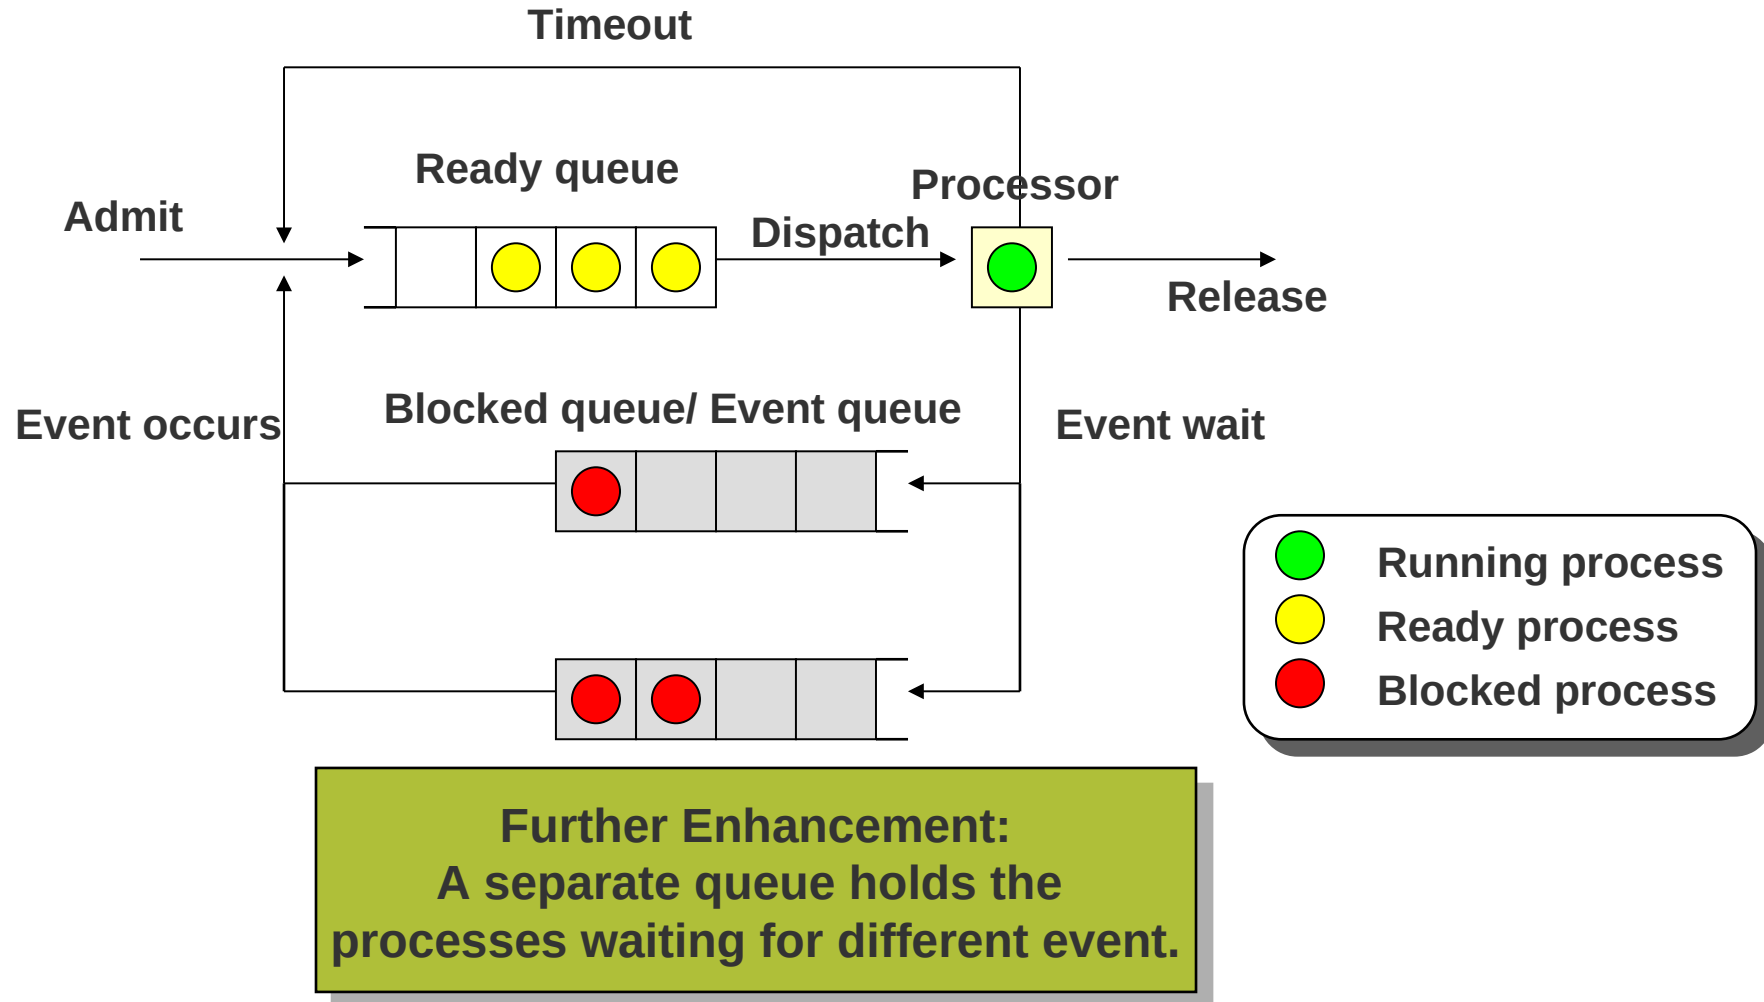
# Five-state Process Model

- Running: currently being run
- Ready: ready to run
- Blocked: waiting for an event (I/O)
- New: just created, not yet admitted to set of run-able processes
- Exit: completed/error exit

# Five-state Process Model



# Blocked Queues

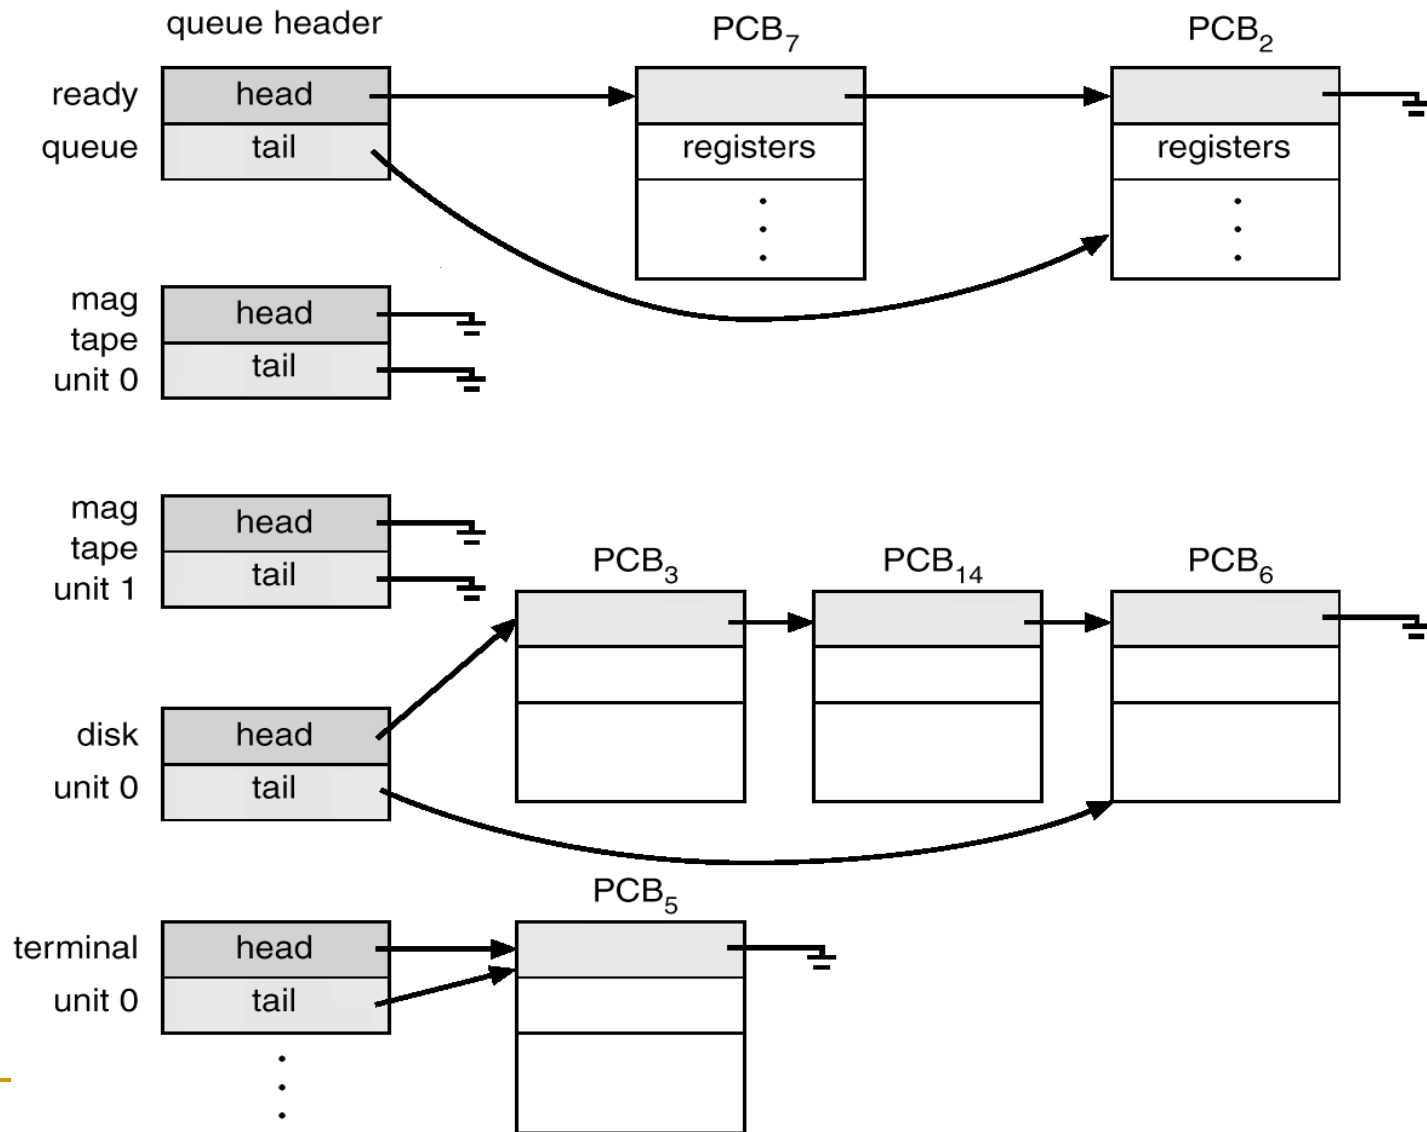




# Scheduling Queues

- The queues are generally stored as linked lists
- A queue header points to the first and the final PCB's in the list
- We extend each PCB to include a pointer field that points to the next PCB in the ready queue

# Scheduling Queues



# Schedulers

- Short term Scheduler or CPU Scheduling
  - Which program is to be run next
- Long term Scheduler or Job Scheduler
  - Which ready jobs should be brought to memory
  - May need to invoke only when a process leaves the system
  - Must make a careful selection

# Process Types

- Most processes can be described as either I/O bound or CPU bound
- I/O bound:
  - Spends more of its time doing I/O than doing computations
- CPU bound:
  - Spends more of its time doing computations than doing I/O
- If all processes are I/O bound,
  - The ready queue will almost always be empty
- If all processes are CPU bound,
  - The I/O waiting queue will almost always be empty, devices will go unused
  - System will again be unbalanced

---

# Process Types

- Sometimes OS may swap a blocked process to disk to free up more memory
- Or to improve process mix
- This is called **Swapping**