

Process Creation

Process Creation

- Includes
 - Build kernel data structures
 - Allocate memory
- Reasons to create a process
 - Submit a new batch job/Start program
 - User logs on to the system
 - OS creates on behalf of a user (printing)
 - Spawned by existing process

Process Termination

- Batch job issues Halt instruction
- User logs off
- Process executes a service request to terminate
- Parent terminates so child processes terminate
- Operating system intervention
 - such as when deadlock occurs
- Error and fault conditions
 - E.g. memory unavailable, protection error, arithmetic error, I/O failure, invalid instruction

Unix Process Creation

- When the system starts up it is running in kernel mode
- There is only one process, the initial process.
- At the end of system initialization, the initial process starts up another kernel process.
- The **init** kernel process has a process identifier of 1.

Process Creation

- These new processes may themselves go on to create new processes.
- All of the processes in the system are descended from the init kernel thread.
- You can see the family relationship between the running processes in a Linux system using the **ps tree** command
- A new process is created by a **fork()** system call

Compiling C++ code

- `g++ test.cpp -o Output`
- Running the code:
- `./Output`

The fork() system call

At the end of the system call there is a new process waiting to run once the scheduler chooses it

- A new data structure is allocated
- The new process is called the child process.
- The existing process is called the parent process.
- The parent gets the child's pid returned to it.
- The child gets 0 returned to it.
- Both parent and child execute at the same point after fork() returns

Unix Process Control

```
int main()
{
    int pid;
    int x = 0;

    x = x + 1;
    fork();
    x = 3;
    printf("%d", x);
}
```


But we want the child process to do something else...

```
int pid;  
int status = 0;  
  
if (pid = fork()) {  
    /* parent */  
    .....  
    pid = wait(&status);  
} else {  
    /* child */  
    .....  
    exit(status);  
}
```

The **fork** syscall returns a zero to the child and the child process ID to the parent

Parent uses **wait** to

sleep until child returns and status.

Fork creates an exact copy of the parent process

Wait variants allow wait on a specific child or

Child process passes status back to parent on **exit**, to report success/failure

Child Process Inherits

- Stack
- Memory
- Environment
- Open file descriptors.
- Current working directory
- Resource limits
- Root directory

Child process DOESNOT Inherit

- Process ID
- Different parent process ID
- Process times
- Own copy of file descriptors
- Resource utilization (initialized to zero)

How can a parent and child process communicate?

- Through any of the normal IPC mechanism schemes.
- But have special ways to communicate
 - For example
 - The variables are replicas
 - The parent receives the exit status of the child

The wait() System Call

- A child program returns a value to the parent, so the parent must arrange to receive that value
- The wait() system call serves this purpose
 - `pid_t wait(int *status)`
 - it puts the parent to sleep waiting for a child's result
 - when a child calls `exit()`, the OS unblocks the parent and returns the value passed by `exit()` as a result of the wait call (along with the pid of the child)
 - if there are no children alive, `wait()` returns immediately
 - also, if there are zombies, `wait()` returns one of the values immediately (and deallocates the zombie)

What is a zombie?

- In the interval between the child terminating and the parent calling `wait()`, the child is said to be a 'zombie'.
- Even though its not running its taking up an entry in the process table.
- The process table has a limited number of entries.

What is a zombie?

- If the parent terminates without calling `wait()`, the child is adopted by `init`.

The solution is:

- Ensure that your parent process calls `wait()` or `waitpid` or etc, for every child process that terminates.

exit()

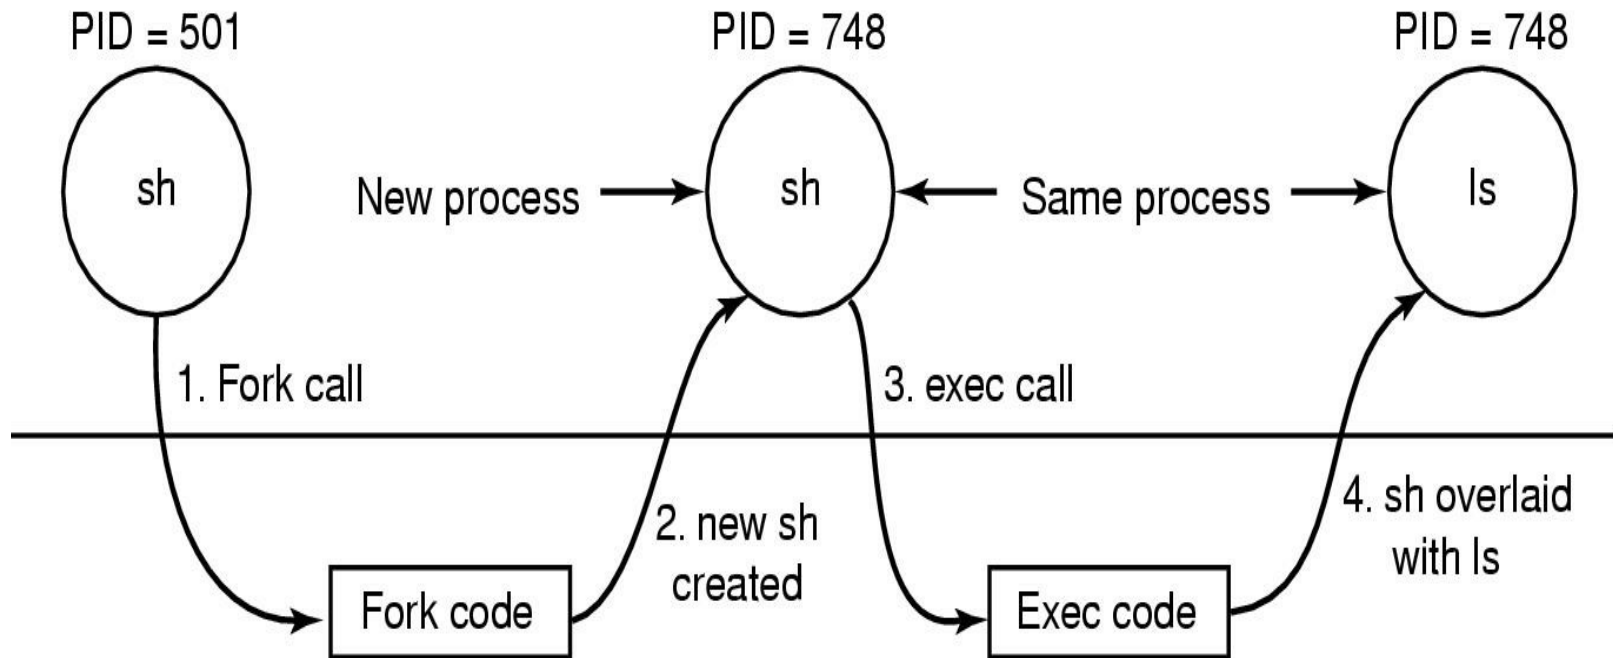
void exit(int *status*);

- After the program finishes execution, it calls *exit()*
- This system call:
 - ❑ takes the “result” of the program as an argument
 - ❑ closes all open files, connections, etc.
 - ❑ deallocates memory
 - ❑ deallocates most of the OS structures supporting the process
 - ❑ checks if parent is alive:
 - If so, it holds the result value until parent requests it, process does not really die, but it enters the zombie/defunct state
 - If not, it deallocates all data structures, the process is dead

execv()

- We usually want the child process to run some other executable
- For Example, *ls*

The *ls* Command



Steps in executing the command *ls* type to the shell

execv

- `int execv(const char *path, char *const argv[]);`
- the current process image with a new process image.
- **path** is the filename to be executed by the child process
- When a C-language program is executed as a result of this call, it is entered as a C-language function call as follows:
 - `int main (int argc, char *argv[]);`
- The **argv** array is terminated by a null pointer.
- The null pointer terminating the **argv** array is not counted in **argc**.