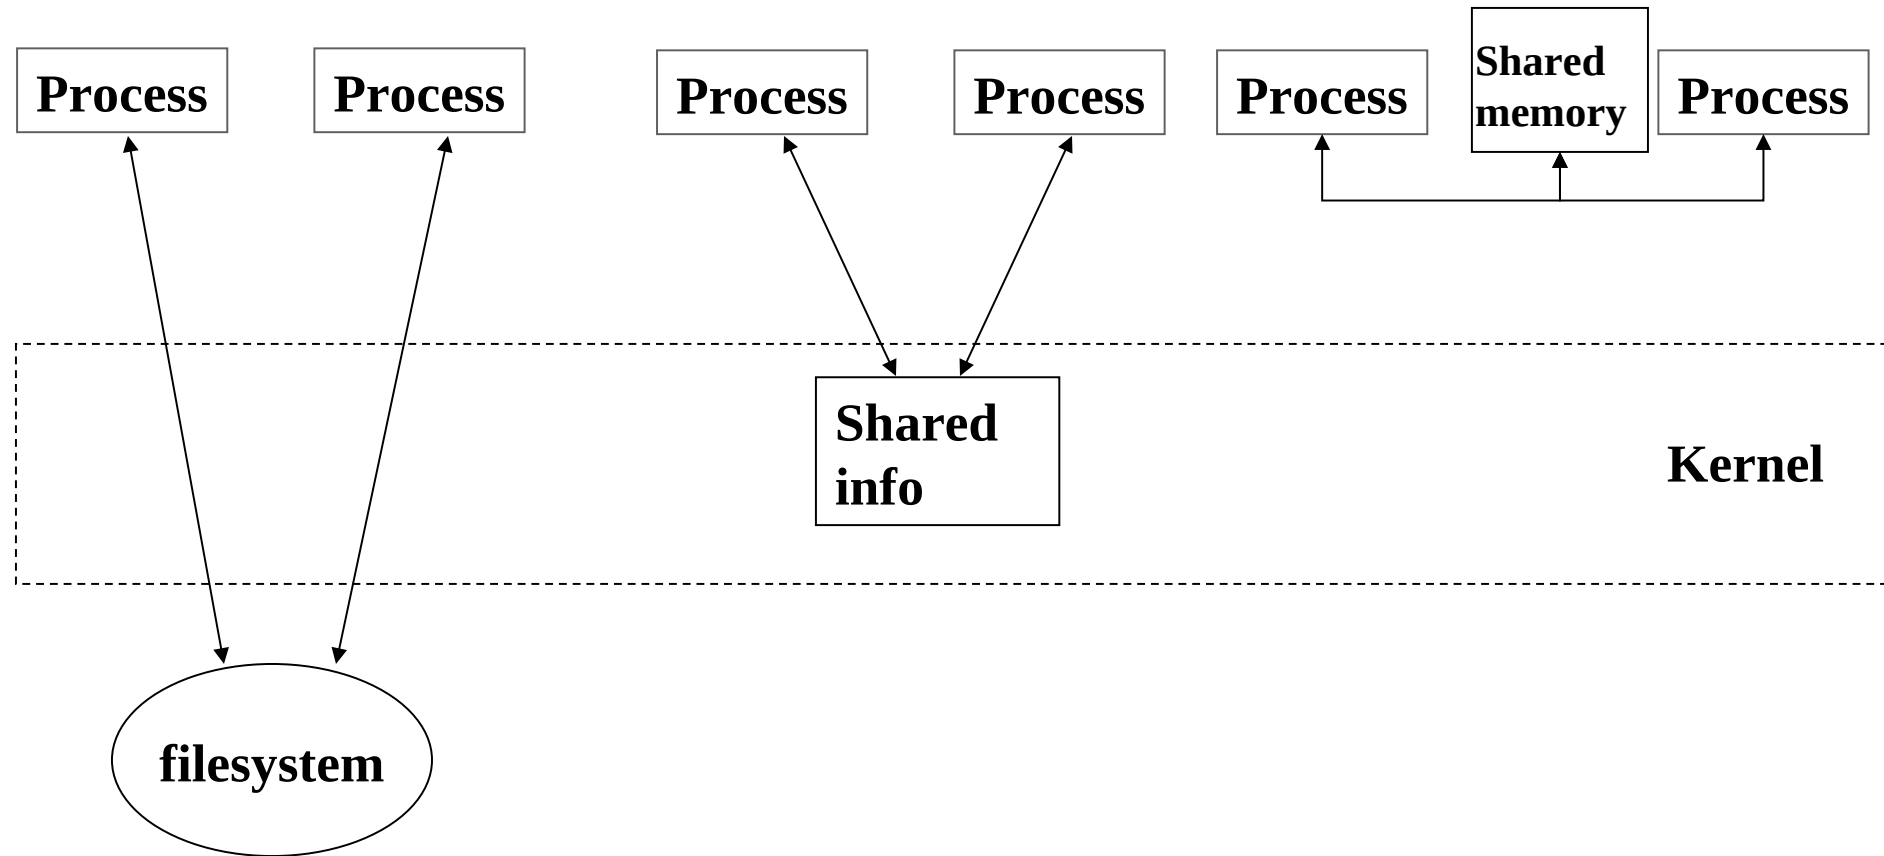


IPC Unix Case Study: Pipes

Unix IPC

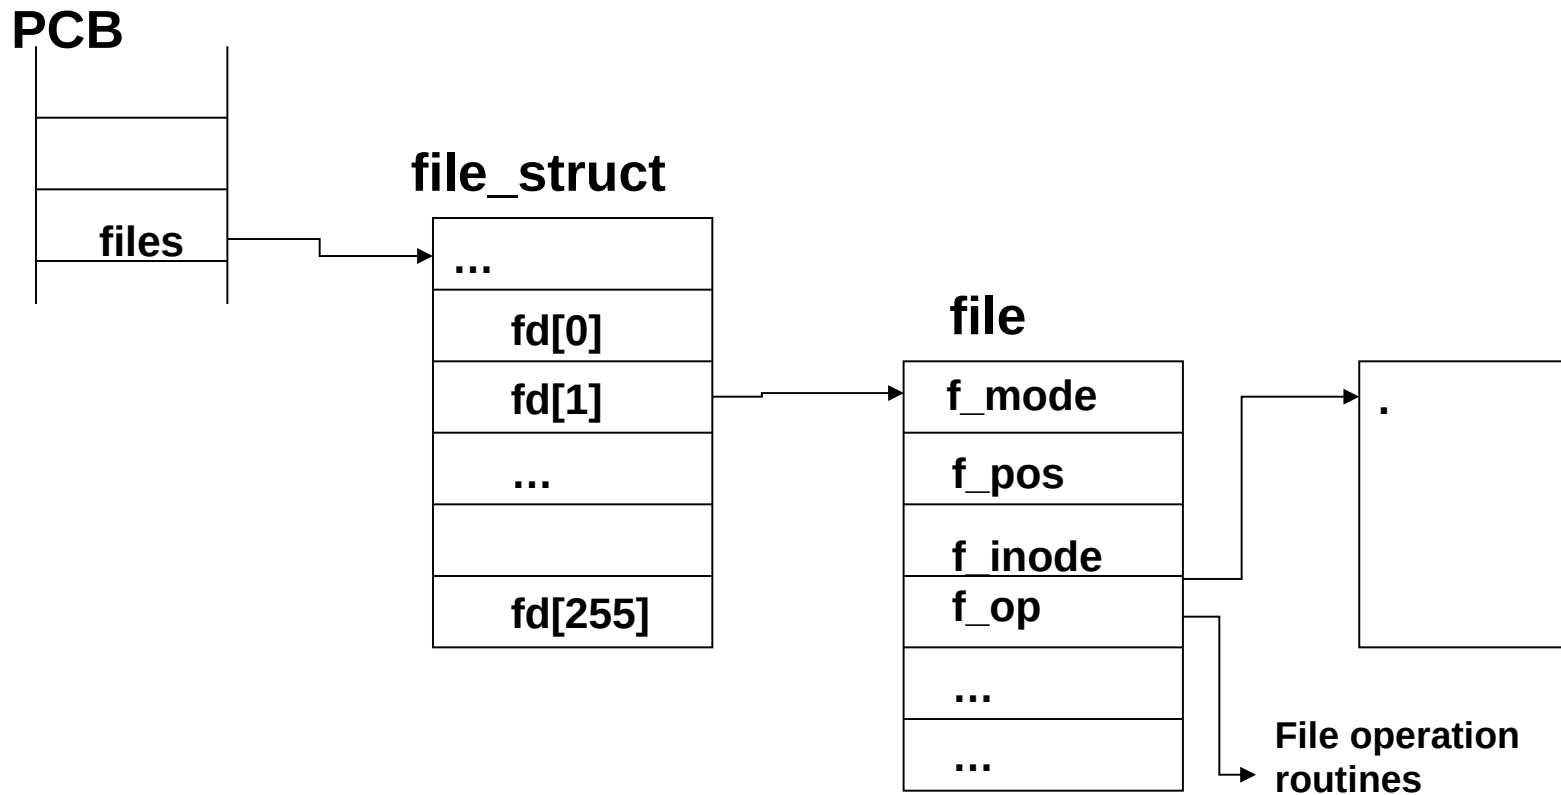


Unix IPC

- Message passing
 - Pipes
 - FIFOs
 - Message queues
- Synchronization
 - Mutexes
 - Condition variables
 - read-write locks
 - Semaphores
- Shared memory
 - Anonymous
 - Named
- Procedure calls
 - RPC
- Event notification
 - Signals

File Descriptors

- The PCB (task_struct) of each process contains a pointer to a file_struct



File Descriptors

- The `files_struct` contains pointers to file data structures
- Each one describes a file being used by this process.
- `f_mode`:
 - describes file mode, read only, read and write or write only.
- `f_pos`:
 - holds the position in the file where the next read or write operation will occur.
- `f_inode`:
 - points at the actual file

File Descriptors

- Every time a file is opened, one of the free file pointers in the `files_struct` is used to point to the new file structure.
- Linux processes expect three file descriptors to be open when they start.
- These are known as standard input, standard output and standard error


File Descriptors

- The program treat them all as files.
- These three are usually inherited from the creating parent process.
- All accesses to files are via standard system calls which pass or return file descriptors.
- standard input, standard output and standard error have file descriptors 0, 1 and 2.

File Descriptors

- `char buffer[10];`
- Read from standard input (by default it is keyboard)
 - `read(0,buffer,5);`
- Write to standard output (by default is is monitor))
 - `write(1,buffer,5);`
- By changing the file descriptors we can write to files
- `fread/fwrite` etc are wrappers around the above `read/write` functions

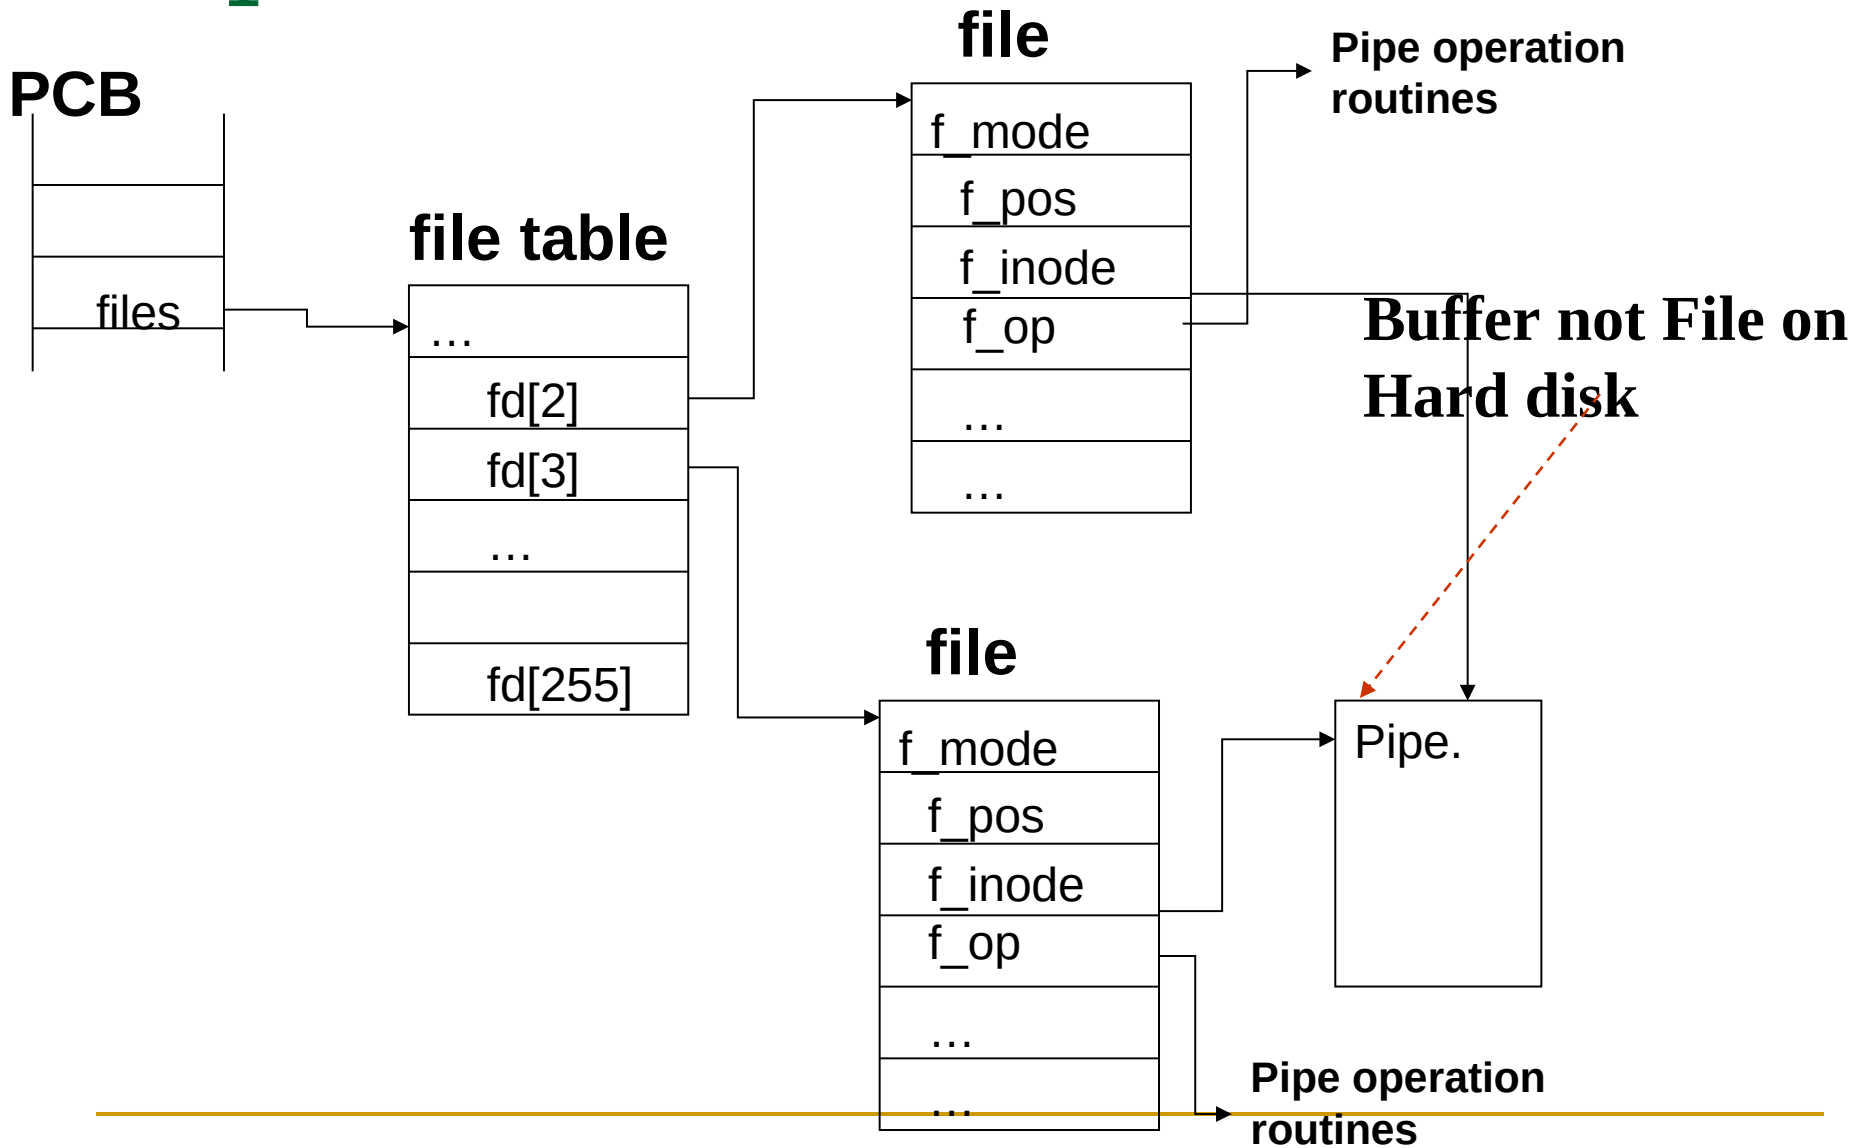
Pipes: Shared info in kernel's memory



Pipes

- A pipe is implemented using two file data structures which both point at the same temporary data node.
- This hides the underlying differences from the generic system calls which read and write to ordinary files.
- Thus, reading/writing to a pipe is similar to reading/writing to a file

Pipes




Pipe Creation

- **#include <unistd.h>**
- **int pipe(int filedes[2]);**
- Creates a pair of file descriptors pointing to a pipe inode
- Places them in the array pointed to by filedes
- filedes[0] is for reading
- filedes[1] is for writing.
- On success, zero is returned.
- On error, -1 is returned

Pipe Creation

```
int main()
{ int pfd[2];

  if (pipe(pfd) == -1)
  { perror("pipe");
    exit(1); }
}
```



Reading/Writing from/to a Pipe

- `int read(int filedescriptor, char *buffer, int bytetoread);`
- `int write(int filedescriptor, char *buffer, int bytetowrite);`

Example

```
int main()
{  int pfd[2];
   char buf[30];
   if (pipe(pfd) == -1) {
       perror("pipe");
       exit(1); }
   printf("writing to file descriptor #%d\n", pfd[1]);
   write(pfd[1], "test", 5);
   printf("reading from file descriptor #%d\n", pfd[0]);
   read(pfd[0], buf, 5);
   printf("read %s\n", buf);
}  write(1, "test", 5);????

read(0, buf, 5);????
```

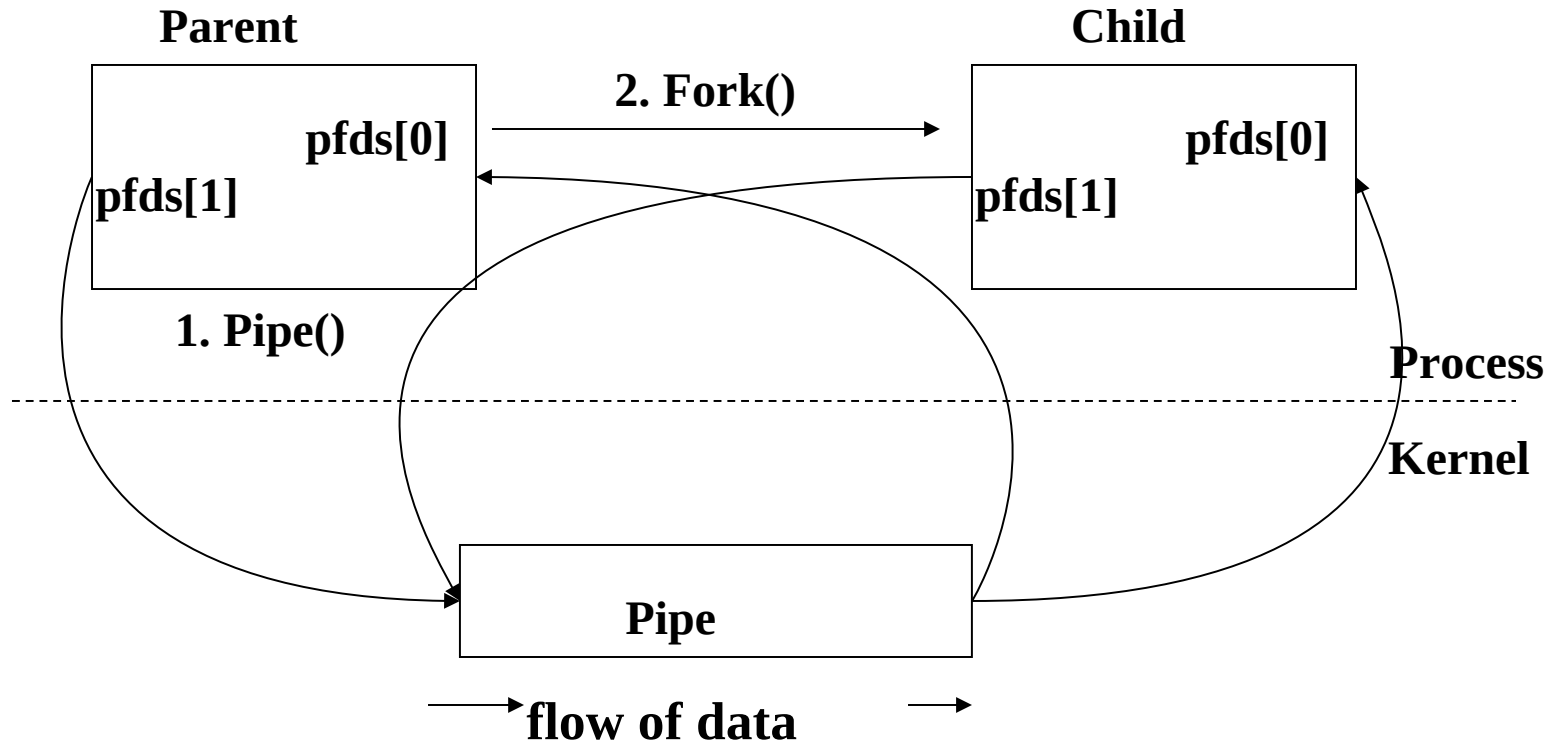
A Channel between two processes

- Remember: the two processes have a parent / child relationship
- The child was created by a `fork()` call that was executed by the parent.
- The child process is an image of the parent process
- Thus, all the **file descriptors** that are opened by the parent are now available in the child.

A Channel between two processes

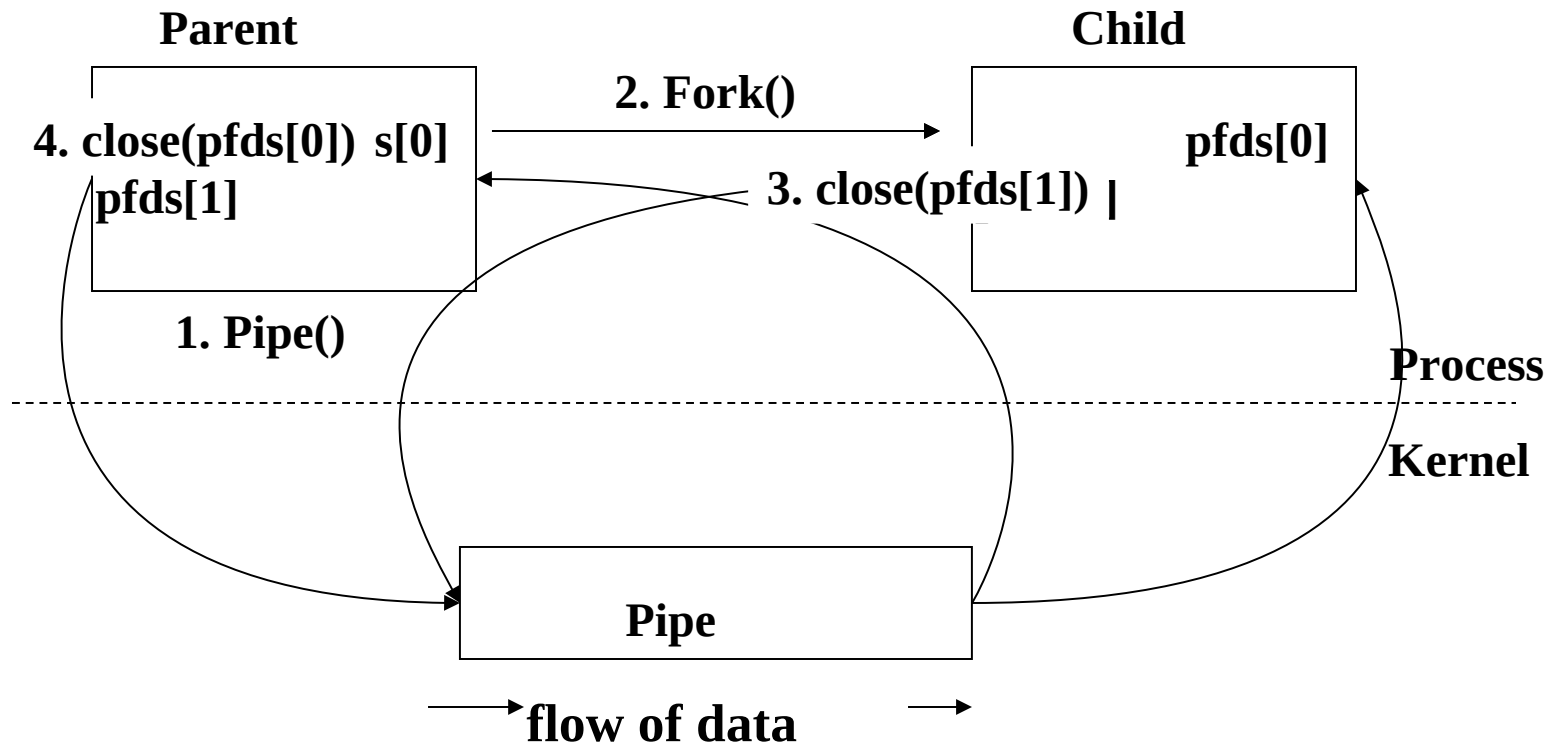
- The **file descriptors** refer to the same I/O entity, in this case a pipe.
- The pipe is inherited by the child
- And may be passed on to the grand-children by the child process or other children by the parent.

A Channel between two processes



A Channel between two processes

- To allow one way communication each process should close one end of the pipe.



Closing the pipe

- The **file descriptors** associated with a pipe can be closed with the `close(fd)` system call
- How would we achieve two way communication?

An Example of pipes with fork

```
int main() {
    int pfd[2];
    char buf[30];
    pipe(pfd); .....1
    if (!fork()) .....2
    { close(pfd[0]); .....3
      printf(" CHILD: writing to the pipe\n");
      write(pfd[1], "test", 5);
      printf(" CHILD: exiting\n");
      exit(0);
    }
    else { close(pfd[1]); .....4
          printf("PARENT: reading from pipe\n");
          read(pfd[0], buf, 5);
          printf("PARENT: read \"%s\"\n", buf);
          wait(NULL);
        }
    }
```